

Distinction assignment lab 3

Jens Madsen, Victor Jarlow

Group 23

To be able to benchmark the code that we produced, we performed some different tweaks to the program:

ParentContinue/Wait:

This tweak describes how the parent should behave after forking its children. If for example the parent has two possible paths it can continue along, it will in the case of the ParentContinue scenario fork only 1 child along one of the paths, and continue itself along the other. In the case of the ParentWait scenario, the parent will fork 2 children, one along each path and then wait from them to join.

ForkAfterSteps/NeighborsDiscovered:

We used two different methods to decide when to fork, either using the amount of steps that a thread as travelled, or the amount of neighbors it has discovered.

ThreadPoolCount:

The ThreadPoolCount describes how many threads the threadpool used to run the program has at its disposal. We used two different levels of ThreadPoolCounts: 2 and 20

The tweaks were tested on 3 separate maps, small map, medium map and large map. The large map was created by us and is an extended version of the medium map. The small and medium maps are the same maps that were supplied to us in the skeleton-code bundle.

The map-sizes in nodes are as follows:

Small: 9X9

Medium: 31X11

Large: 31X21


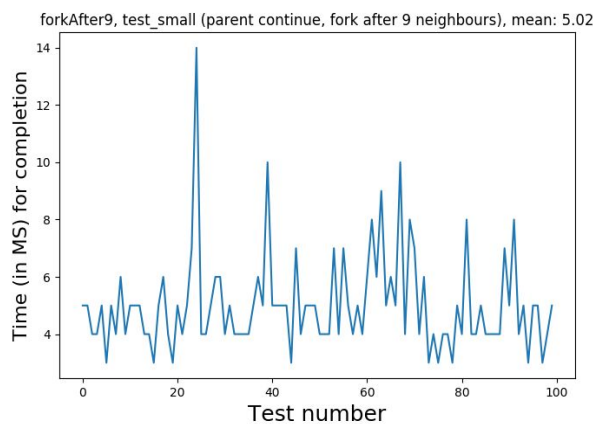
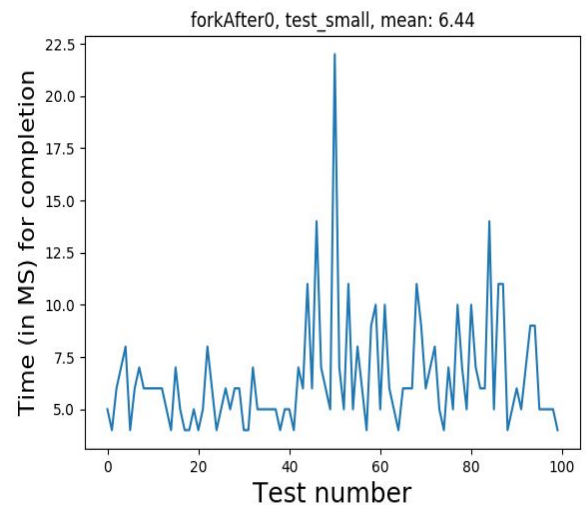
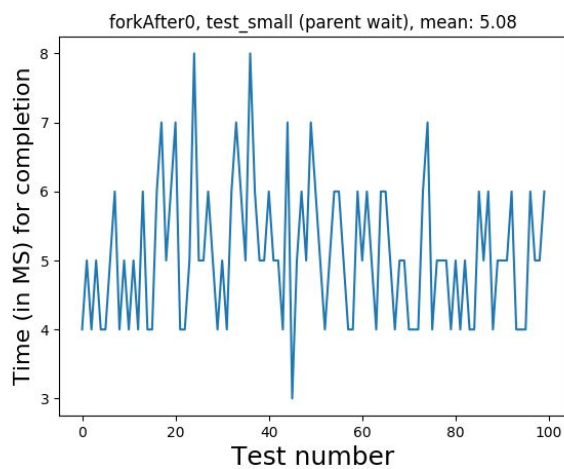
Very Large: 119X572

Results





Initially we used 3 different map sizes: small, medium and large. We noticed that the maps were too small to see any significant differences between the tweaks that we applied in each case. Therefore we decided to make a significantly larger map (very large), which we hoped would highlight any differences and make them more visible.

Since the differences between tweaks in the map sizes small, medium and large were so small we decided that the results of those benchmarks were inconclusive, but we still decided to enclose them in this report.


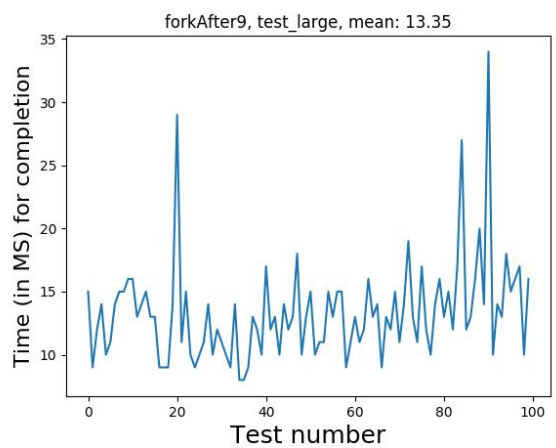
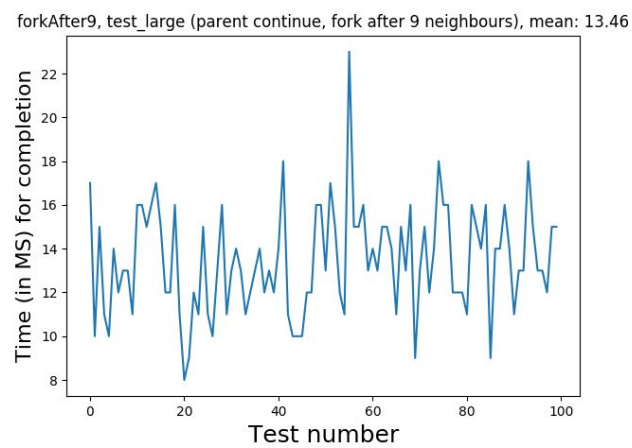

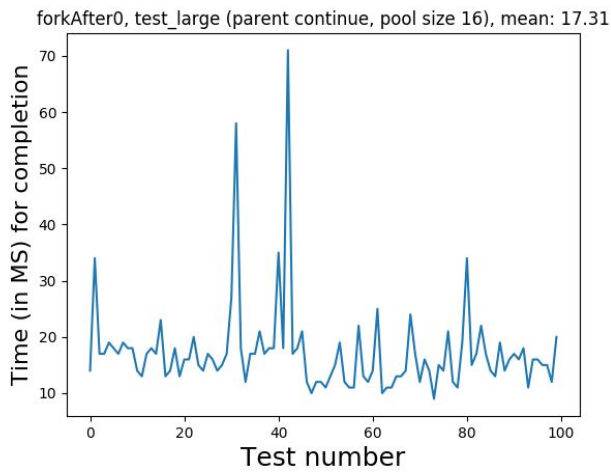


Small-sized Map



Medium-sized Map



Large-sized Map



Very Large-sized Map




Figure shows the total completion times for all tests on the very large map.




Figure shows the mean completion time for the tests on the very large map.

	continue	wait
Fork after 1000000 step (never forks)	143	143
Fork after 0 step	1496	5893
Fork after 0 neighbor	1474	6567
Fork after 200 step	2424	1542
Fork after 200 neighbor	2198	1264
Fork after 200 pool size 2	N/A	2300
Fork after 200 pool size 20	N/A	1426

Table 1. Benchmarking times using the very large map.

Conclusions

In this chapter we will discuss what conclusions can be drawn from the benchmarking test. As previously mentioned we will be focusing on the benchmarking tests made using the very large map, since the results obtained using the other maps did not show enough variance between tests to rule out randomness.

When we made the map large (verylarge.map) we found that forkAfter a high number of iterations was much more efficient than forking after a few iterations. We think this is because the map is large and wide (and we put the goal the furthest away) which leads to a depth-first search being more effective than a breadth-first search. Using this map and goal position is the worst case for a breadth-first search. Generating new threads for each of the elements in the frontier mimics a breadth-first search which leads to forking new threads all the time and is less efficient than not forking at all (which mimics a depth-first search) for this type of map.

We noticed that if the threads are set to fork after 0 steps/neighbors and the tweak where the parent has to wait for the children is used, the time it takes to find the goal increases dramatically.

The same is true for the inverse, when threads are set to fork after 200 steps/neighbors and the tweak that parents continue after forking children.

We don't understand why this phenomenon occurs, but perhaps we need to conduct more tests to be able to come to a conclusion.